
e516: Engineering Cloud Computing

Indiana University

Gregor von Laszewski

2018-05-10

Contents

1 e516: Engineering Cloud Computing	2
1.1 Course Description	2
1.2 Course Objectives	3
1.3 Learning Outcomes	3
1.4 Syllabus	3
1.5 Assessment	6
1.6 Assignments	7
1.7 Calendar	8
1.8 Incomplete	9
1.9 Other classes I523, I524, B649, E516, E616	9

1 e516: Engineering Cloud Computing

- Lecture Notes: [ePub](#)
- Indiana University
- Faculty: Dr. Gregor von Laszewski (laszewski@gmail.com)
- Credits: 3
- Hardware: You will need a computer to take this class, a phone, tablet, or chrome book is not sufficient.
- Prerequisite(s): Knowledge of a programming language, the ability to pick up other programming languages as needed, willingness to enhance your knowledge from online resources and additional literature. You will need access to a *modern* computer that allows using virtual machines and/or containers. If such a system is not available to you can also use cloud vms we provide and if you opt to do so one or more Raspberry's PI. All residential students will have access to a total of 200 Raspberry PIs. Online students can opt to purchase one or more based on our materials list that we will release throughout the semester. All students will have access to a cloud.
- Course Description: [Link](#)

This is an introductory class. In case you like to do research and more advanced topics, consider taking an independent study with Dr. von Laszewski.

1.1 Course Description

This course covers basic concepts on programming models and tools of cloud computing to support data intensive science applications. Students will get to know the latest research topics of cloud plat-

forms, parallel algorithms, storage and high level language for proficiency with a complex ecosystem of tools that span many disciplines.

1.2 Course Objectives

The course has the following objectives:

- Provide a basic introduction to cloud computing
- Introduce the concept of cloud data centers
- Get familiar with cloud infrastructure as a Service such as OpenStack, Azure, or AWS
- Get familiar with cloud infrastructure such as Docker and Kubernetes
- Program cloud services
- Understand the differences between virtual machines and containers
- Develop sophisticated programming language independent REST services
- Learn advanced programming models for clouds such as Map/Reduce, Messaging, and GraphQL
- Exploration of Go for cloud computing
- Demonstrate knowledge of clouds while developing a significant project
- Explore state-of-the-art cloud technologies and services while providing a section and summary and commenting on its use for the cloud
- Learn how edge computing is enhancing cloud services and infrastructure
- Learn how to set up a cloud based on using commodity hardware

1.3 Learning Outcomes

- Be able to explain the concepts of the cloud computing paradigm including its paradigm shift, its characteristics, and the advantages. Contrast them with the challenges and disadvantages.
- Be able to identify infrastructure and programming models needed to support real world applications.
- Be able to implement a real world application or deploy a cloud and its services.
- Be able to conduct sophisticated performance analysis of cloud services.
- Be able to communicate the results through tutorials, manual, and reports.
- Be able to work in a team to develop collaboratively software or contribute collaboratively to develop sections explaining how to use clouds.

1.4 Syllabus

The topics are subject to change.

Dates	Unit	Title	Description
Week 1	1	Introduction	Gregor von Laszewski Class summary Definition of Cloud Computing
Week 2	2	Tools	Tools and Services - Virtual Box - Vagrant - Github - Linux
Week 3	3	Python	Python - Introduction - Installation - Interactive Python - Editors - Basic Language Features - Modules - Data Management - Matplotlib - Cloudmesh Commandshell CMD5 - OpenCV - Secchi Disk
Week 4	4		Data Center
Week 5	5	Architectures	- NIST Big Data Reference Architecture - Cloud Architectures - NIST Big Data Reference Architecture
Week 6	6	Virtualization	Virtualization, Qemu, KVM, Virtual machines - Qemu
Week 7	7	Infrastructure	Infrastructure as a Service

Dates	Unit	Title	Description
			- Azure ↔
			- AWS ↔
			- OpenStack ↔
		Chameleon Cloud	- Chameleon Cloud
			- Resources ↔
			- Hardware ↔
			- Charge ↔
			- Quick start ↔
			- KVM user guide ↔
			- CLI ↔
			- Horizon ↔
			- Heat ↔
			- Baremetal ↔
			- FAQ ↔
Week 8	8	Virtualization II	Containers, Docker, Kubernetes
Week 9	9	Programming	Python for Cloud Computing
			- Libcloud ↔
			- Github as a Cloud Service ↔
			- REST Services ↔
			- Rest Services with OpenAPI ↔
			- OpenAPI Spec ↔
			- OpenAPI Codegen ↔
Week 10	10	Map/Reduce	Map/Reduce, Hadoop, Spark, and others
Week 11	11	Messaging	Messaging
Week 12	12	Messaging	- MQTT ↔
			- GraphQL
Week 13	13	Go	Go for the Cloud

Dates	Unit	Title	Description
Week 14,15	14	Edge Computing	Edge Computing and the Cloud
Week 16	15	Project Selection	Project selection for inclusion in the Proceedings

Students need only to do one project. The project is conducted throughout the entire semester.

- Example chapters are indicated with \leftrightarrow
- Dates may change as the semester evolves
- The project is a long term assignment (and are ideally worked on weekly by residential students). It is the major part of the course grade.
- Sections and chapters prepare you for documenting a technical aspect related to cloud computing. It is a preparation for a document that explains how to execute your project in a reproducible manner to others.
- Additional lectures will be added that allow easy management of the project. These lectures can be taken any time when needed.

1.5 Assessment

This course is focusing on the principal “Learning by Doing” which is assessed by simple graded and non-graded activities. The assessment may include comprehension of the material taught, programming assignments, participation in online discussion forums, or the contribution of additional material to the class showcasing your comprehension.

The comprehension is also measured by the development of a chapter and sections in markdown that can be distributed and replicated to other students. This is done in preparation for the project that must include a simple deployment and runtime instruction set.

The main deliverable of the class is a project. The project is assessed through the following artifacts:

1. Deployment and install instructions,
2. Project report (typically 2-3 pages per month, tutorial and chapters can be reused if possible),
3. Working project code that can be installed and executed in reproducible manner by a third party
4. Code developed by the project team distributed in github.com
5. Project progress notes checked into github throughout the semester. Each week the project progress is reported and will be integrated into the final grade.

6. Discussions or progress reports with the instructors can be conducted on online and residential class hours.

The grade distribution is as follows

- 10% Comprehension Activities
- 10% Sections
- 10% Chapter
- 70% Project

As the project is the main deliverable of the course it is obvious that those starting a week before the deadline will not succeed in this class. The project will take a significant amount of time and fosters the principal of “Learning by Doing” at all stages throughout the semester.

The class will not have a regular midterm, but it is expected that you have worked on your project and can provide a snapshot of the progress outlining the goals of the project and how you will achieve these goals till the end of the semester.

The final Project is due 3 weeks before the semester ends. Issues with your project ought to have been discussed before this deadline with the TA's. The TAs will in the next 3 weeks go over the projects and evaluate major and minor issues that you may be able to fix without penalty. Larger changes will receive a grade penalty. The last fix (upon approval) possible will be 2 weeks before semester end without penalty. It is recommended that you are available till the last day of class.

1.6 Assignments

Dates	Unit	Title	Description
due Week 8	A1	Sections	Contribute significant sections. Do not develop redundant or duplicated content.
due Week 8	A2	Chapter	Contribute a significant chapter that may use your section to the class documentation. Do not develop redundant or duplicated content.
Due Week 13	A3	Project Type A	Build a cloud out of Raspberry Pis Kubernetes, Hadoop, SLURM + OpenAPI Service,
		Project Type B	Build a Significant OpenAPI REST Service
		Project Type C	Build an Edge Service Interfacing with a Cloud
		Project Type D	Contribute to the new Cloudmesh code

Dates	Unit	Title	Description
		Project Type E	Your own Project Type A, B, C, D [upon approval)

1.7 Calendar

Assignment #	Event		Date
	Full Term		16 Weeks
	<i>Begins</i>		Week 1
1, 2	Bio, Notebook	assigned	Week 1
1, 2	Bio, Notebook	due	Week 2
3	Sections	assigned	Week 3
4	Chapter	assigned	Week 3
5	Project	selection or proposal	Week 4
5	Project	Draft	Week 8
3	Section	due	Week 10
4	Chapter	due	Week 10
5	Project	due	Week 12
5	Project (no penalty)	improvements	Week 13
5	Project (with penalty)	improvements	Week 14
	Final Deliverables due		Week 14
	<i>Grading</i>		Week 15, 16
	<i>Ends</i>		Week 16

- TA's must be available till all grades have been submitted.
- Bio: a formal 3 paragraph Bio
- Notebook: a markdown in which you record your progress of this class in bullet form
- All times are in EST
- Dependent on class progress Comprehension Assignments may be added

1.8 Incomplete

Please see the university regulations for getting an incomplete. However, as this class uses state-of-the-art technology that changes frequently, you must expect that an incomplete may result in significant additional work on your behalf as your project may need significant updates on infrastructure, technology, or even programming models used. It is best to complete the course within one semester.

1.9 Other classes I523, I524, B649, E516, E616

The classes CSCI B-649 I524 are the same as E516 but are offered within different departments:

- CSCI B-649 Cloud Computing is the same as E516 but for computer science students.
- I524 is the same as E516 but for Data Engineering Students

The class E616 is a bit different as it assumes you have taken E516 as prerequisite or know its content including REST, virtual machines, containers, and Hadoop. In this class you are expected to focus on advanced concepts that build upon 516. This includes DevOps, Julia, and any topic of 516 with enhancements. The focus will allow you to deliver contributions based on your knowledge as part of sections, a chapter you propose to add to the lecture notes, and a more intensive project. It also allows PhD students to identify topics that can be leading together with Dr. von Laszewski as a co-author to a workshop or conference paper.

Dr. von Laszewski also offers independent studies on specific research topics. The result of such independent studies is a high quality report or paper.

It is possible to take E516 and E616 or an independent study at the same time. The result will be a more intense project.

I523 is a different class and focusses on Big Data applications and use cases. You do not have to take I523 as prerequisite if you like to take E516.